

01 PROBLEM UNDERSTANDING AND OBJECTIVE

Clinical Document Intelligence Hub · 2-day prototype

Clinical and administrative staff read the same documents twice.

Intake forms, discharge summaries, lab reports and physician notes arrive as text, PDF, scans and faxes. Someone reads each one for the same handful of facts, then re-keys them elsewhere. Slow, inconsistent between readers, and the errors it produces are silent.

The hard part is not reading the document. It is whether a clinician can *check* what was read — quickly, per field, without re-reading the chart.

OBJECTIVE

Ingest one unstructured clinical document and return a decision-ready patient summary card:

- structured fields, each with a confidence and its source quote
- a risk flag a clinician can defend
- concrete recommended next steps

THREE CONSTRAINTS THAT SHAPED EVERY DECISION

Absent data must stay absent

Silent imputation is the failure mode that actually hurts someone. Missing fields become the literal string “not documented”; if three or more NEWS2 parameters are missing the band is INDETERMINATE, not a score built on partial vitals.

The risk score must not be generated


A prose rationale from a model cannot be tested, audited, or held still while a prompt is edited. NEWS2 is ordinary Python, covered by 18 tests.

Every value must be traceable

Each field carries a verbatim source quote, expandable in the UI. Verifying a dose should take a click, not a re-read.

02 SOLUTION ARCHITECTURE AND DESIGN FLOW

One model call, one schema, one deterministic scorer



the typed record forks

DETERMINISTIC src/triage.py · src/compare.py

NEWS2 risk score — plain Python

RCP scoring from the extracted vitals; every point attributable to a reading. A second document for the same patient is diffed by the same scorers, so “improved” always means fewer NEWS2 points. 46 tests. The model never sees this step.

GENERATIVE summary · reasoning · next steps


Narrative, grounded in quotes

Prose a clinician reads in under 20 seconds, plus recommended actions ordered most urgent first and an explicit list of what is missing from the document.

The model extracts; Python decides. The risk band cannot drift when a prompt is edited, a clinician can be shown which reading contributed which point, and the whole scoring layer is testable without an API key.

03 IMPLEMENTATION HIGHLIGHTS

The card a reviewer sees, and the call that produces it



Risk band, then the fields it rests on. Each value carries a confidence marker and an expandable **source** holding the verbatim quote it came from.

MEASURED ON THE SIX BUNDLED SAMPLES

95

source quotes audited
against their documents

0

containing content
absent from the source

59

tests — scoring, diffing
and failure handling

~29 s

per document,
end to end

\$0.00

free tier —
reproducible by anyone

Model: gemini-3.6-flash · extract.py is the only module aware of the provider.

ONE CALL — THE API ENFORCES THE SCHEMA

```
interaction = client.interactions.create(
    model=model,
    system_instruction=SYSTEM_PROMPT,
    input=_input_blocks(document),
    response_format={
        "type": "text",
        "mime_type": "application/json",
        "schema": response_schema(),
    },
    store=False, # nothing retained
)
```

No JSON parsing, no repair loop, no prose-to-struct step that can drift.
`response_schema()` inlines the `$defs/$ref` Gemini does not document.

04 CHALLENGES AND LEARNINGS

The confidence score was decoration — measuring it is what proved that

The prototype reported **high** confidence on almost every field. That looks like success until you ask whether the signal *moves*. So I rendered the same ED triage note at three degradation levels and compared confidence against ground truth.

BEFORE confidence is flat while accuracy collapses

Degradation	Vitals right	Confidence
As shipped	6 / 6	21 high, 1 med
Degraded	5 / 6	23 high, 0 med
Severe	3 / 6	20 high, 0 med

AFTER the signal tracks legibility

Degradation	Vitals right	Confidence
As shipped	6 / 6	8 high, 16 med
Degraded	4 / 6	6 high, 18 med
Severe	1 / 6	0 high, 17 med

At severe degradation the model read temperature as 37.9 °C instead of 38.9, and heart rate as 118 instead of 124 — and labelled both **high** confidence. The fix was conceptual, not mechanical: confidence had been written as a statement of **belief**. It is now a statement about the **source** — any digit that could be a different digit under this much blur is **medium**, however sure the model feels. Clean digital text is unaffected: 105 high / 0 medium.

ALSO LEARNED

The cheap model was a trap

flash-lite: 4x faster, same NEWS2, but it fabricated 3 of 65 source quotes — one carried the reference range from the row above. Rejected.

A green tick is a claim

Uniform high confidence tells a clinician there is nothing to check. That is worse than showing no confidence at all.

One document cannot check itself

Diffing two encounters exposed a triage bug that had survived 18 tests and five samples: “nil allergies, confirmed” was being flagged as a missing allergy history.

Degradation levels generated reproducibly (fixed seed) by scripts/make_scanned_sample.py. · The after-fix ladder reproduced exactly (8 → 6 → 0 high) on a second independent run.

05 DEMO SUMMARY AND NEXT STEPS

Six synthetic documents, two modalities, one paired encounter

END-TO-END RESULTS AGAINST THE LIVE API

Document	Input	Classified as	NEWS2	Risk
ED triage note — suspected urosepsis	text	triage_note	17	HIGH
same note, printed → photocopied → faxed	image	triage_note	17	HIGH
Discharge summary — resolved sepsis	text	discharge_summary	0	ROUTINE
New-patient intake form — headaches	text	intake_form	0	ROUTINE
same patient, 7 days later — discharge	text	discharge_summary	0	ROUTINE
Lab report — potassium 6.8, no vitals recorded	text	lab_report	—	INDETERMINATE

A degraded fax produced the identical decision — **NEWS2 17, HIGH, all 7 parameters, all 7 labs, the same 5 flags.** No OCR stage.

Load the same patient's discharge summary and the app diffs the encounters: **NEWS2 17 → 0, HIGH → ROUTINE, lactate critical → normal, five flags resolved.**

GIVEN MORE TIME

- **Close the audit trail.** One run produced a vitals quote that skipped a line while reading as continuous. Mark elisions; give each vital its own quote.
- **Widen the comparison.** It diffs two documents; a real record is a timeline. Identity matching also needs more than MRN-then-surname.
- **Pin the non-determinism.** Values are stable across runs; formatting is not (BD vs twice daily). Normalise in the schema.

TO TAKE THIS BEYOND A PROTOTYPE

- **Gate on patient category.** NEWS2 is validated for acutely ill adults — not paediatric or obstetric patients, and the tool does not detect those cases.
- **Human-in-the-loop by default.** Reviewed-and-accepted should be a recorded state with reviewer and timestamp, not an assumption.
- **Structured interoperability.** Emit FHIR resources rather than bespoke JSON, so output can re-enter the record it came from.